

hw3\_submission.md

# Homework 3

Name: Clemence Goh (1002075)

Talked to: Cheryl Goh

## Question 1

1 a)

Objective is to find a rank-one matrix which is closest to the incomplete matrix M, where:

$$M = \begin{bmatrix} 0 & 1 & x \\ 1 & x & 1 \\ x & 1 & 2 \end{bmatrix}$$

Code:

```
import numpy as np

# using alternating least squares method for collaborative filtering:
# http://www.quuxlabs.com/blog/2010/09/matrix-factorization-a-simple-tutorial-and-implementation-in-python/
def get_error(Q, X, Y, W):
    return np.sum((W * (Q - np.dot(X, Y)))**2)

def matrix_factorization(R, P, Q, K, steps=10000, alpha=0.0002, beta=0.02):
    Q = Q.T
    for step in range(steps):
        for i in range(len(R)):
            for j in range(len(R[i])):
                if R[i][j] > 0:
                    eij = R[i][j] - np.dot(P[i, :], Q[:, j])
                    for k in range(K):
                        P[i][k] = P[i][k] + alpha * (2 * eij * Q[k][j] - beta * P[i][k])
                        Q[k][j] = Q[k][j] + alpha * (2 * eij * P[i][k] - beta * Q[k][j])
                    eR = np.dot(P, Q)
                    e = 0
            for j in range(len(R[i])):
                if R[i][j] > 0:
                    e = e + pow(R[i][j] - np.dot(P[i, :], Q[:, j]), 2)
                    for k in range(K):
                        e = e + (beta/2) * (pow(P[i][k], 2) + pow(Q[k][j], 2))
            if e < 0.0001:
                break
    return P, Q.T

def initialiseVariables():
    R = np.array(
        [[0, 1, np.nan],
         [1, np.nan, 1],
         [np.nan, 1, 2]]
    )
```

```

    )
    N = len(R)
    M = len(R[0])
    K = 1 # number of rows

    P = np.random.rand(N, K)
    Q = np.random.rand(M, K)
    return M, N, K, P, Q, R

def checkAgainst(_nR):
    desired_matrix = np.array(
        [[0, 1, np.nan],
         [1, np.nan, 1],
         [np.nan, 1, 2]]
    )
    for i in range(len(desired_matrix)):
        for j in range(len(desired_matrix[0])):
            # skip these numbers
            if i == 0 and j == 2 or i == 1 and j == 1 or i == 2 and j == 0:
                continue
            else:
                if round(_nR[i][j]) != desired_matrix[i][j]:
                    return False
    return True

if __name__ == '__main__':
    M, N, K, P, Q, R = initialiseVariables()

    nP, nQ = matrix_factorization(R, P, Q, K)
    nR = np.dot(nP, nQ.T)
    print(nR)
    print(nP)
    print(nQ)
    print(checkAgainst(nR))

```

Results:

```

A:
[[1.02515492 0.82214388 1.35229942]
 [0.85242192 0.68361713 1.12444436]
 [1.43113499 1.14772787 1.88783469]]


U:
[[0.97078532]
 [0.8072133 ]
 [1.35523404]]

V:
[[1.05600579]
 [0.84688536]
 [1.39299533]]

```

## Question 2

2 a)



Data plotted:

Code:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn import svm

def showGraph(_plt):
    _plt.show()

def plotGraphData():
    csv = 'kernel.csv'
    data = np.genfromtxt(csv, delimiter=',')

    X = data[:, 1:]
    Y = data[:, 0]
    plt.scatter(X[:, 0], X[:, 1], c=Y)
    return X, Y, plt

def part1():
    _, _, plt = plotGraphData()
    showGraph(plt)

# function to train SVM based on scikit package
# set gamma to 0.5, kernel to rbf
def trainSVM():
    # part a
    clf = svm.SVC(gamma=0.5, kernel='rbf')
    X, Y, plt = plotGraphData()
    clf.fit(X, Y)
    return clf, X, Y
```

Result:

```
clf:
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=0.5, kernel='rbf',
```

```
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

2 b)


Code:

```
# this is for part b
def decision(_x1, _x2, _clf):
    x = np.array([[_x1, _x2]])
    return _clf.decision_function(x)[0]
```

2 c)

Code:


```
# part c
def visualiseClassifier():
    clf, X, Y = trainSVM()
    vdecision = np.vectorize(decision, excluded=[2])
    x1list = np.linspace(-8.0, 8.0, 100)
    x2list = np.linspace(-8.0, 8.0, 100)
    X1, X2 = np.meshgrid(x1list, x2list)
    Z = vdecision(X1, X2, clf)
    cp = plt.contourf(X1, X2, Z)
    plt.colorbar(cp)
    plt.scatter(X[:, 0], X[:, 1], c=Y, cmap='gray')
    plt.show()
```



Graph Plotted:

## Question 3

3.



Data matrix X:

3 a)

code:

```
import numpy as np
from numpy.linalg import norm
import matplotlib.pyplot as plt
from scipy.optimize import fmin_l_bfgs_b as minimize
from Homework.Week6.utils import normalize, tile_raster_images, sigmoid, ravelParameters, unravelParameters
from Homework.Week6.utils import initializeParameters, computeNumericalGradient

def initialiseVariables():
    nV = 8 * 8 # number of visible units
    nH = 25 # number of hidden units
    dW = 0.0001 # weight decay term
    sW = 3 # sparsity penalty term
    npy = 'images.npy'
    X = normalize(np.load(npy))
    return nV, nH, dW, sW, X

# part a
def sparseAutoencoderCost(theta, nV, nH, dW, sW, X):
    W1, W2, b1, b2 = unravelParameters(theta, nH, nV)
    n = X.shape[0]
    z2 = np.dot(X, W1) + np.dot(np.ones((n, 1)), b1.T)
    a2 = sigmoid(z2)
    z3 = np.dot(a2, W2) + np.dot(np.ones((n, 1)), b2.T)
    a3 = sigmoid(z3)
    eps = a3 - X
    loss = 0.5 * np.sum(eps**2) / n
    decay = 0.5 * (np.sum(W1**2) + np.sum(W2**2))

    # Compute sparsity terms and total cost
    rho = 0.01
    a2mean = np.mean(a2, axis=0).reshape(nH, 1)
    k1 = np.sum(rho * np.log(rho / a2mean) + (1 - rho) * np.log((1 - rho) / (1 - a2mean)))
    dk1 = -rho / a2mean + (1 - rho) / (1 - a2mean)
    cost = loss + dW * decay + sW * k1
    d3 = eps * a3 * (1 - a3)
```

```

d2 = (sw*dk1.T+np.dot(d3,w2.T))*a2*(1-a2)
W1grad = np.dot(X.T, d2)/n + dw * W1
W2grad = np.dot(a2.T, d3)/n + dw * W2
b1grad = np.dot(d2.T, np.ones((n, 1)))/n
b2grad = np.dot(d3.T, np.ones((n, 1)))/n
grad = ravelParameters(W1grad,W2grad,b1grad,b2grad)
print(' .',end="")
return cost,grad

def checkCostGradFunction():
    nV, nH, dW, sW, X = initialiseVariables()
    theta = initializeParameters(nH, nV)
    cost, grad = sparseAutoencoderCost(theta, nV, nH, dW, sW, X)
    return cost, grad

```

Running checkCostGradFunction() gives:

```

.
Cost: 54.352568640271166

Grad: [ 0.76502459  0.94329594  0.79515806 ...  0.02468031  0.00289279
        -0.02947588]

```

3 b)

code:

```

def compareGradients():
    # from previous
    nV, nH, dW, sW, X = initialiseVariables()
    theta = initializeParameters(nH, nV)
    cost, grad = sparseAutoencoderCost(theta, nV, nH, dW, sW, X)

    print('\nComparing numerical gradient with backprop gradient')
    num_coords = 5
    indices = np.random.choice(theta.size, num_coords, replace=False)
    numgrad = computeNumericalGradient(lambda t: sparseAutoencoderCost(t, nV, nH, dW, sW, X)[0], theta, indices)
    subnumgrad = numgrad[indices]
    subgrad = grad[indices]
    diff = norm(subnumgrad - subgrad) / norm(subnumgrad + subgrad)
    print('\n', np.array([subnumgrad, subgrad]).T)
    print('The relative difference is', diff)

```

Running comapreGradients() gives Result:

```


Comparing numerical gradient with backprop gradient
. . . . .
[[-0.03965281 -0.03965281]
 [ 0.72180184  0.72180184]
 [ 0.85801614  0.85801614]
 [-0.00907039 -0.00907039]
 [ 0.89547885  0.89547885]]
The relative difference is 4.9672839287460054e-11

```

3 c)

Code:

```
def trainNeuralNetwork():  
    nV, nH, dW, sW, X = initialiseVariables()  
    theta = initializeParameters(nH, nV)  
    opttheta, cost, messages = minimize(sparseAutoencoderCost, theta, fprime=None, maxiter=400,  
                                       args=(nV, nH, dW, sW, X))  
    W1, W2, b1, b2 = unravelParameters(opttheta, nH, nV)  
    plt.imshow(tile_raster_images(X=W1.T, img_shape=(8, 8), tile_shape=(5, 5), tile_spacing=(1, 1)), cmap='gray')  
    plt.show()
```



Output Graph:

## Question 4

---



4a) Given: Likelihood  $\Rightarrow P(S|\theta) = \prod_{w \in W} \theta_w^{n(w)}$   
 Training Loss  $\Rightarrow \mathcal{L}_n(\theta) = \frac{1}{n} \sum_{w \in W} -n(w) \log \theta_w$

where:  $\theta_w \geq 0$  for all  $w \in W$   
 $\sum_{w \in W} \theta_w = 1$

using Lagrange:  $\mathcal{L}(\theta, \lambda) = \mathcal{L}_n(\theta) + \lambda (\sum_{w \in W} \theta_w - 1)$   
 $= \frac{1}{n} \sum_{w \in W} -n(w) \log \theta_w + \lambda (\sum_{w \in W} \theta_w - 1)$

$$\frac{\partial \mathcal{L}(\theta, \lambda)}{\partial \theta_w} = -\frac{1}{n} \frac{n(w)}{\theta_w} + \lambda \stackrel{\text{set}}{=} 0$$

multiply across by  $\theta_w$ :

$$\theta_w \lambda = \frac{1}{n} n(w)$$

$$\lambda \sum \theta_w = \frac{1}{n} \sum n(w) \quad \text{from slides, given } \sum_{w \in W} \theta_w = 1$$


$$\therefore \lambda = \frac{1}{n} \sum n(w) \quad \text{--- (1)}$$

sub  $\lambda$  into  $\theta_w \lambda = \frac{1}{n} n(w)$ :  
 $\Rightarrow w \rightarrow w'$  for  $\lambda$

$$\theta_w \frac{1}{n} \sum n(w') = \frac{1}{n} n(w)$$

$$\hat{\theta}_w = \frac{n(w)}{\sum_{w' \in W} n(w')} \quad \text{(shown)} *$$





4 b) Gaussian:  
Given  $L_n(\mu, \sigma^2) = \frac{d}{2} \log(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{x \in S} \|x - \hat{\mu}\|^2$ :

For  $\hat{\mu}$ :  $\frac{\partial L}{\partial \mu} = \frac{1}{2\pi\sigma^2} \sum_{x \in S} -2(x - \hat{\mu}) \stackrel{\text{set}}{=} 0$

multiply across by  $-\pi\sigma^2$ :

$$\sum_{x \in S} (x - \hat{\mu}) = 0$$

Distribute summation:

$$\sum_{x \in S} x - n\hat{\mu} = 0$$

$$\therefore \hat{\mu} = \frac{1}{n} \sum_{x \in S} x \quad * \text{ (shown)}$$

For  $\sigma^2$ : let  $\sigma^2$  be  $\theta$

$$\Rightarrow \frac{\partial L}{\partial \theta} = \frac{d}{2} \frac{1}{2\pi\theta} \cdot 2\pi + \frac{1}{2n} \cdot -\frac{1}{\theta^2} \cdot \sum_{x \in S} \|x - \hat{\mu}\|^2 = 0 \quad (\text{set as } 0).$$

$$\frac{d}{2\theta} - \frac{1}{2n\theta^2} \sum_{x \in S} \|x - \hat{\mu}\|^2 = 0$$

multiply across by  ~~$\theta^2$~~   $2\theta^2$ :

$$d\theta - \frac{1}{2n} \sum_{x \in S} \|x - \hat{\mu}\|^2 = 0$$

$\rightarrow$  sub  $\theta = \sigma^2$ :

$$d\sigma^2 = \frac{1}{2n} \sum_{x \in S} \|x - \hat{\mu}\|^2$$

$$\sigma^2 = \frac{1}{2n} \sum_{x \in S} \|x - \hat{\mu}\|^2 \quad * \text{ (shown)}$$

