

Classification and Regression

Sunday, October 28, 2018 2:42 PM

2017 and 2016:

1.

Point Loss $\mathcal{L}_1(\theta, \theta_0; x, y)$	Predictor	Technique	Algorithm*	Learning Cost
$\mathcal{L}_S(y - (\theta^\top x + \theta_0))$	$f(x; \theta, \theta_0) = \theta^\top x + \theta_0$	Linear Regression	Exact Solution, Gradient Descent	$\mathcal{L}_n(\theta, \theta_0) = \frac{1}{n} \sum_{\text{data}(x, y)} \frac{1}{2} (y - (\theta^\top x + \theta_0))^2$
$\mathcal{L}_S(y - (\theta^\top x + \theta_0)) + \frac{\lambda}{2} \ \theta\ ^2$	$f(x; \theta, \theta_0) = \theta^\top x + \theta_0$	Ridge Regression	Exact Solution, Gradient Descent	$\mathcal{L}_n(\theta, \theta_0) = \frac{1}{n} \sum_{\text{data}(x, y)} \frac{1}{2} (y - (\theta^\top x + \theta_0))^2 + \frac{\lambda}{2} \ \theta\ ^2$
$\mathcal{L}_H(y(\theta^\top x + \theta_0))$	$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$	Linear Classification using Hinge Loss	Gradient Descent	$\mathcal{L}_n(\theta, \theta_0) = \frac{1}{n} \sum_{\text{data}(x, y)} \max\{1 - y(\theta^\top x + \theta_0), 0\}$
$\mathcal{L}_H(y(\theta^\top x + \theta_0)) + \frac{\lambda}{2} \ \theta\ ^2$	$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$	Support Vector Machine with Slack Variables	Gradient Descent	$\mathcal{L}_n(\theta, \theta_0) = \frac{1}{n} \sum_{\text{data}(x, y)} \max\{1 - y(\theta^\top x + \theta_0), 0\} + \frac{\lambda}{2} \ \theta\ ^2$
$\mathcal{L}_Z(y(\theta^\top x + \theta_0))$	$h(x; \theta, \theta_0) = \text{sign}(\theta^\top x + \theta_0)$	Perceptron (with Offset)	Perceptron Algorithm	$\mathcal{L}_n(\theta, \theta_0) = \frac{1}{n} \sum_{\text{data}(x, y)} \mathbb{I}[y(\theta^\top x + \theta_0) < 1]$
$\mathcal{L}_L(y(\theta^\top x + \theta_0))$	$p(y x, \theta, \theta_0) = \text{sigmoid}(y(\theta^\top x + \theta_0))$	Logistic Regression	Gradient Descent	$\mathcal{L}_n(\theta, \theta_0) = \frac{1}{n} \sum_{\text{data}(x, y)} -\log(p(y x, \theta, \theta_0))$

2.

Gradient for Linear Regression without the offset θ_0 :	$\nabla_{\theta} \mathcal{L}_n(\theta; \mathcal{S}_n) = \frac{1}{n} \sum_{(x, y) \in \mathcal{S}_n} x(\theta^\top x - y)$
2015: Gradient for Logistic Regression without the [Ridge Regression] Ridge regression optimizes what objective function?	$\nabla_{\theta} \mathcal{L}_n(\theta; \mathcal{S}_n) = \frac{1}{n} \sum_{(x, y) \in \mathcal{S}_n} x(\text{sigmoid}(\theta^\top x) - \mathbb{I}[y = 1])$

=>

mean squared error with squared ℓ_2 -norm $\|w\|_2^2 = \sum_{d=1}^D w_d^2$ on weights w

[Regression Error]

Why is the following measure no good objective function for measuring the error in a regression problem? The error is computed between ground truth y_i and prediction $f(x_i)$ as given by the function

$$E = \frac{1}{N} \sum_{i=1}^N (f(x_i) - y_i)^3$$

Hint: you can imagine what can happen if this objective is used with a linear model: $f(x_i) = w^\top x_i$.

Ans: The loss function z^3 is negative for negative errors z . Since the objective is to minimize the loss, the trained parameters will end up favoring large negative errors.

[Logistic Regression]

1. We consider here a discriminative approach for solving the classification problem illustrated in Figure 1

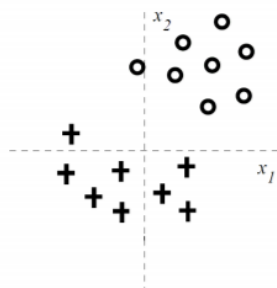


Figure 1: The 2-dimensional labeled training set, where '+' corresponds to class $y=1$ and 'O' corresponds to class $y=0$.

- a. We attempt to solve the binary classification task depicted in Figure 1 with the simple linear logistic regression model



Notice that the training data can be separated with zero training error with a linear separator. Consider training regularized linear logistic regression models where we try to maximize

$$\sum_{i=1}^n \log(P(y_i|x_i, w_0, w_1, w_2)) - Cw_j^2$$

for very large C . The regularization penalties used in penalized conditional log-likelihood estimation are $-Cw_j^2$, where $j = \{0, 1, 2\}$. In other words, only one of the parameters is regularized in each case. Given the training data in Figure 1, how does the training error change with regularization of each parameter w_j ? State whether the training error increases or stays the same (zero) for each w_j for very large C . Provide a brief justification for each of your answers.

- i. By regularizing w_2 :

Ans: Increases. When we regularize w_2 , the resulting boundary can rely less and less on the value of x_2 and therefore becomes more vertical. For very large C , the training error increases as there is no good linear vertical separator of the training data

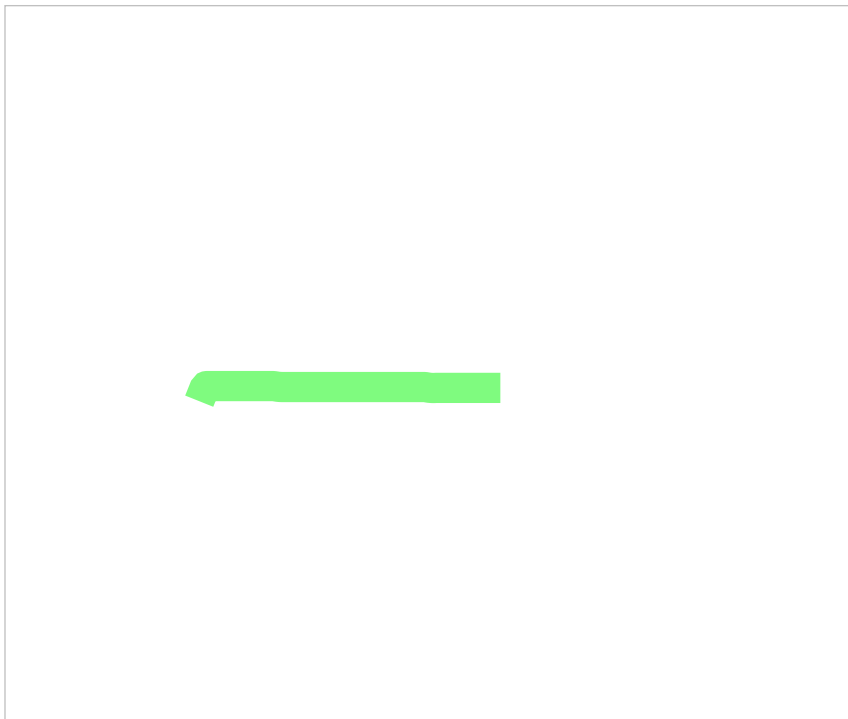
- ii. By regularizing w_1 :

Ans: Remains the same. When we regularize w_1 , the resulting boundary can rely less and less on the value of x_1 and therefore becomes more horizontal and the training data can be separated with zero training error with a horizontal linear separator.

- iii. By regularizing w_0 :

Ans: Increases. When we regularize w_0 , then the boundary will eventually go through the origin (bias term set to zero). Based on the figure, we can not find a linear boundary through the origin with zero error. The best we can get is one error.

b.



L1 regularization => Lasso Regression

Lasso Regression (Least Abs
“absolute value of magnitude”
function.

$$\sum_{i=1}^n$$

L2 regularization => Ridge Regression

c.

[3 pts] For very large C , with the same L1-norm regularization for w_1 and w_2 as above, which value(s) do you expect w_0 to take? Explain briefly. (Note that the number of points from each class is the same.) (You can give a range of values for w_0 if you deem necessary).

SOLUTION: For very large C , we argued that both w_1 and w_2 will go to zero. Note that when $w_1 = w_2 = 0$ the log-probability of labels becomes a finite value, which is

that when $w_1 = w_2 = 0$, the log-probability of labels becomes a finite value, which is equal to $n \log(0.5)$, i.e. $w_0 = 0$. In other words, $P(y = 1|\vec{x}, \vec{w}) = P(y = 0|\vec{x}, \vec{w}) = 0.5$. We expect so because the number of elements in each class is the same and so we would like to predict each one with the same probability, and $w_0 = 0$ makes $P(y = 1|\vec{x}, \vec{w}) = 0.5$.

d.

[3 pts] Assume that we obtain more data points from the '+' class that corresponds to $y=1$ so that the class labels become unbalanced. Again for very large C , with the same L1-norm regularization for w_1 and w_2 as above, which value(s) do you expect w_0 to take? Explain briefly. (You can give a range of values for w_0 if you deem necessary).

SOLUTION: For very large C , we argued that both w_1 and w_2 will go to zero. With unbalanced classes where the number of '+' labels are greater than that of 'o' labels, we want to have $P(y = 1|\vec{x}, \vec{w}) > P(y = 0|\vec{x}, \vec{w})$. For that to happen the value of w_0 should be greater than zero which makes $P(y = 1|\vec{x}, \vec{w}) > 0.5$.